

SOURCE: README.md

Legacy Application Modernization Platform - Setup

? Quick Start

After reinstalling Docker Desktop, you have **two deployment options**:

Option 1: PowerShell Script (Recommended)

```
# Full deployment with all services and AI models
.\rebuild-platform.ps1

# Skip AI model downloads (if already downloaded)
.\rebuild-platform.ps1 -SkipModels

# Skip monitoring stack
.\rebuild-platform.ps1 -SkipMonitoring

# Validation only
.\rebuild-platform.ps1 -Validate
```

Option 2: Docker Compose

```
# Start all services
docker-compose up -d

# View logs
docker-compose logs -f

# Stop all services
docker-compose down

# Stop and remove volumes (CAUTION: data loss)
docker-compose down -v
```

? --

? Prerequisites

- ? **Docker Desktop** installed and running
- ? **Windows 11 Pro** with WSL2 enabled
- ? **GPU Support** configured (for Ollama)
- ? **PowerShell 5.1+ or PowerShell Core 7+**

- ? **Minimum 32GB RAM** (64GB+ recommended)
- ? **500GB+ available disk space** (for models and data)

? --

?? Platform Architecture

Infrastructure Services

Service	Port	Purpose	Credentials
PostgreSQL	5432	Relational database	'modernization' / 'modernization'
Redis	6379	Caching layer	No authentication
MinIO	9000, 9001	Object storage	'minioadmin' / 'minioadmin123'
ChromaDB	8000	Vector database	No authentication
Ollama	11434	AI model server	No authentication

Monitoring Stack

Service	Port	Purpose	Credentials
Prometheus	9090	Metrics collection	No authentication
Grafana	3000	Visualization	'admin' / 'admin123'
Loki	3100	Log aggregation	No authentication
Jaeger	16686	Distributed tracing	No authentication

AI Models (Ollama)

- ? **DeepSeek Coder 33B** ('deepseek-coder:33b-instruct-q4_K_M') - ~18GB
- ? **Qwen2.5-Coder 32B** ('qwen2.5-coder:32b-instruct-q4_K_M') - ~18GB

? --

? Validation & Health Checks

Quick Health Check

```
.\\validate-platform.ps1
```

Detailed Health Check

```
.\\validate-platform.ps1 -Detailed
```

Continuous Monitoring

```
# Check every 30 seconds  
.\\validate-platform.ps1 -Continuous -Interval 30
```

? --

?? Manual AI Model Management

Pull Additional Models

```
# List available models  
docker exec ollama ollama list  
  
# Pull a specific model  
docker exec ollama ollama pull deepseek-coder:6.7b-instruct-q4_K_M  
  
# Pull larger models (70B+)  
docker exec ollama ollama pull qwen2.5-coder:72b-instruct-q4_K_M
```

Test Model Inference

```
# Test DeepSeek Coder  
docker exec ollama ollama run deepseek-coder:33b-instruct-q4_K_M "Write a Python fu  
  
# Test Qwen2.5-Coder  
docker exec ollama ollama run qwen2.5-coder:32b-instruct-q4_K_M "Explain what a mic
```

? --

? Accessing Services

Web Interfaces

```
? **MinIO Console**: http://localhost:9001  
? **Prometheus**: http://localhost:9090  
? **Grafana**: http://localhost:3000  
? **Jaeger**: http://localhost:16686
```

API Endpoints

```
? **Ollama API**: http://localhost:11434  
? **ChromaDB API**: http://localhost:8000  
? **Loki API**: http://localhost:3100
```

Database Connections

? *PostgreSQL*:

```
# Via Docker
docker exec -it postgres psql -U modernization -d legacy_modernization

# Connection string
postgresql://modernization:modernization123@localhost:5432/legacy_modernization
```

? *Redis*:

```
# Via Docker
docker exec -it redis redis-cli

# Connection string
redis://localhost:6379
```

? --

? Troubleshooting

Check Container Status

```
# View all containers
docker ps -a

# View specific container logs
docker logs -f <container-name>

# Check container resource usage
docker stats
```

Common Issues

1. Port Already in Use

```
# Find process using port
netstat -ano | findstr :<port>

# Kill process
Stop-Process -Id <PID> -Force
```

2. Container Won't Start

```
# Check logs
docker logs <container-name>

# Restart container
docker restart <container-name>

# Remove and recreate
docker rm -f <container-name>
.\rebuild-platform.ps1
```

3. Ollama GPU Not Detected

```
# Verify GPU support in Docker
docker run --rm --gpus all nvidia/cuda:12.2.0-base-ubuntu22.04 nvidia-smi

# Check Ollama container GPU access
docker exec ollama nvidia-smi
```

4. Out of Memory

```
# Check memory usage
docker stats --no-stream

# Restart Docker Desktop
# Or reduce number of running models

? --
```

?? Cleanup & Reset

Stop All Containers

```
docker stop $(docker ps -aq)
```

Remove All Containers

```
docker rm $(docker ps -aq)
```

Remove Volumes (?? DATA LOSS)

```
docker volume rm postgres-data redis-data minio-data chromadb-data ollama-data prom
```

Complete Reset

```
# Stop and remove everything
docker stop $(docker ps -aq)
```

```
docker rm $(docker ps -aq)
docker volume prune -f
docker network prune -f

# Rebuild from scratch
.\rebuild-platform.ps1
```

? --

? Monitoring & Metrics

Prometheus Targets

Visit <http://localhost:9090/targets> to verify all scrape targets are UP.

Grafana Dashboards

1. Log in to Grafana: <http://localhost:3000> (admin / admin123)
2. Add Prometheus data source: - URL: ‘<http://prometheus:9090>’
3. Add Loki data source: - URL: ‘<http://loki:3100>’
4. Import dashboards for: - Ollama metrics - System metrics - Container metrics - GPU metrics

Common Queries

? *Prometheus**:

```
# Ollama request rate
rate(ollama_request_total[5m])

# Container CPU usage
rate(container_cpu_usage_seconds_total[5m])

# Memory usage
container_memory_usage_bytes / 1024 / 1024 / 1024
```

? --

? Security Considerations

Default Credentials (Change in Production!)

- ? PostgreSQL: ‘modernization’ / ‘modernization123’
- ? MinIO: ‘minioadmin’ / ‘minioadmin123’
- ? Grafana: ‘admin’ / ‘admin123’

Network Isolation

- ? Services are isolated in Docker networks
- ? Only necessary ports are exposed to host
- ? Consider using reverse proxy (Traefik/Nginx) for production

Data Persistence

- ? All data stored in Docker volumes
- ? Volumes persist after container removal
- ? Regular backup recommended for production

? --

? Next Steps

1. **Test Core Workflows**: - Legacy code analysis - Architecture diagram generation - Jira ticket generation
2. **Configure Monitoring**: - Set up Grafana dashboards - Configure alert rules in Prometheus
- Set up log forwarding to Loki
3. **Deploy Application Layer**: - REST API services - Web UI components - Integration services
4. **Scale to Cloud**: - Plan Kubernetes migration - Set up CI/CD pipelines - Configure auto-scaling

? --

? Support & Resources

Platform Status

```
.\\validate-platform.ps1 -Detailed
```

Docker Desktop Settings

- ? **WSL Integration**: Enabled
- ? **GPU Support**: Enabled
- ? **Memory**: 64GB+ allocated
- ? **Disk**: 500GB+ available

Useful Links

- ? [Docker Documentation](<https://docs.docker.com/>)
- ? [Ollama Documentation](<https://github.com/ollama/ollama>)
- ? [Prometheus Documentation](<https://prometheus.io/docs/>)
- ? [Grafana Documentation](<https://grafana.com/docs/>)

? --

? License & Attribution

Legacy Application Modernization Platform Copyright ? 2026 - Cloud, AI & DevOps Architecture Team

? --

- ? *Last Updated**: 2026-01-16
- ? *Platform Version**: 1.0.0
- ? *Deployment Environment**: Windows 11 Pro + Docker Desktop + WSL2

SOURCE: INSTALL.md

INSTALLATION INSTRUCTIONS

? Download Complete!

You have downloaded the **Legacy Application Modernization Platform** deployment package.

? --

? Package Contents

This ZIP file contains: - 'rebuild-platform.ps1' - Main deployment script - 'validate-platform.ps1' - Health check and validation script - 'docker-compose.yml' - Docker Compose configuration - 'prometheus.yml' - Prometheus monitoring configuration - 'loki-config.yml' -

Loki log aggregation configuration - ‘promtail-config.yml’ - Promtail log collector configuration - ‘README.md’ - Complete documentation - ‘QUICK-REFERENCE.md’ - Command reference guide - ‘TROUBLESHOOTING.md’ - Troubleshooting guide

? --

? Quick Start (3 Steps)

Step 1: Extract Files

Extract the ZIP file to a folder on your desktop, for example: ““
C:\LegacyModernizationPlatform\““

Step 2: Open PowerShell

Right-click on the folder and select **"Open in Terminal"** or: ““powershell # Navigate to the folder cd C:\LegacyModernizationPlatform ““

Step 3: Run Deployment

```
# Deploy everything  
.\\rebuild-platform.ps1
```

That's it! The script will: ? Deploy all infrastructure services ? Deploy monitoring stack ? Download AI models ? Validate everything

? --

?? Expected Deployment Time

? **Infrastructure Services**: 2-3 minutes
? **Monitoring Stack**: 1-2 minutes
? **AI Models Download**: 15-30 minutes (one-time)
? **Total**: ~20-35 minutes for first run

Subsequent runs (with models already downloaded): ““powershell .\\rebuild-platform.ps1 -SkipModels ““ Takes only: ~3-5 minutes

? --

? Prerequisites Check

Before running, ensure:

- ? Docker Desktop is installed and running
- ? WSL2 is enabled
- ? GPU support is configured (for Ollama)
- ? At least 32GB RAM available
- ? At least 100GB free disk space
- ? PowerShell 5.1+ or PowerShell Core 7+

Check Docker: “powershell docker --version docker info “

? --

? Script Parameters

```
# Full deployment (first time)
.\rebuild-platform.ps1

# Skip AI model downloads (if already downloaded)
.\rebuild-platform.ps1 -SkipModels

# Skip monitoring stack (minimal deployment)
.\rebuild-platform.ps1 -SkipMonitoring

# Validation only (check if everything is running)
.\rebuild-platform.ps1 -Validate
```

? --

? Post-Deployment Validation

After deployment completes, validate everything: “powershell .\validate-platform.ps1 -Detailed “

This shows:

- Service health status
- Port availability
- Container status
- AI models installed
- Connectivity tests

? --

? Access Your Platform

After deployment, access services at:

? *Infrastructure:**
? PostgreSQL: ‘localhost:5432’
? Redis: ‘localhost:6379’
? MinIO Console: http://localhost:9001
? ChromaDB: http://localhost:8000
? Ollama: http://localhost:11434

? *Monitoring:**
? Prometheus: http://localhost:9090
? Grafana: http://localhost:3000 (admin/admin123)
? Loki: http://localhost:3100
? Jaeger: http://localhost:16686

? --

? Need Help?

1. **Check README.md** - Complete documentation 2. **Check QUICK-REFERENCE.md** - Common commands 3. **Check TROUBLESHOOTING.md** - Fix common issues

? *Common Issues:**

Port already in use: “powershell netstat -ano | findstr :<port> Stop-Process -Id <PID> -Force”

Container won’t start: “powershell docker logs <container-name> ”

Complete reset: “powershell .\rebuild-platform.ps1 ”

? --

? Verify Deployment Success

You should see output like: “ ? Platform Deployment Completed Successfully! ?

Infrastructure Services: ? PostgreSQL: localhost:5432 ? ? Redis: localhost:6379
? ? MinIO Console: http://localhost:9001 ? ? ChromaDB: http://localhost:8000 ? ?
Ollama: http://localhost:11434 ?

Monitoring Stack: ? Prometheus: http://localhost:9090 ? ? Grafana: http://localhost:3000 ? ? Loki: http://localhost:3100 ? ? Jaeger: http://localhost:16686 ? ?

? --

Support

If you encounter any issues:

1. Run validation: “powershell .\validate-platform.ps1 -Detailed”
2. Check container logs: “powershell docker logs <container-name>”
3. Review TROUBLESHOOTING.md for solutions
4. Generate diagnostic report: “powershell docker info > diagnostic-report.txt docker ps -a >> diagnostic-report.txt .\validate-platform.ps1 -Detailed >> diagnostic-report.txt”

? --

Next Steps

After successful deployment:

1. **Test Core Services** “powershell # Test Ollama docker exec ollama ollama list

Test PostgreSQL docker exec postgres psql -U modernization -d legacy_modernization -c "SELECT version();"
- # Test Redis docker exec redis redis-cli PING “
2. **Configure Grafana Dashboards** - Visit http://localhost:3000 - Add Prometheus data source: http://prometheus:9090 - Add Loki data source: http://loki:3100
3. **Run Your Workflows** - Legacy code analysis - Architecture diagram generation - Jira ticket generation

? --

? Default Credentials

Remember to change these in production:

? **PostgreSQL**: modernization / modernization123

? **MinIO**: minioadmin / minioadmin123

? **Grafana**: admin / admin123

? --

? Enjoy Your Platform!

You now have a complete AI-powered Legacy Application Modernization Platform running locally with full monitoring and observability!

For detailed documentation, refer to README.md

? *Last Updated*: 2026-01-16

? *Platform Version*: 1.0.0

SOURCE: QUICK-REFERENCE.md

QUICK REFERENCE CARD - Common Commands

? DEPLOYMENT COMMANDS

Full platform rebuild

```
.\rebuild-platform.ps1
```

Rebuild without AI models

```
.\rebuild-platform.ps1 -SkipModels
```

Rebuild without monitoring

```
.\rebuild-platform.ps1 -SkipMonitoring
```

Docker Compose deployment

```
docker-compose up -d
```

Docker Compose with rebuild

```
docker-compose up -d --build --force-recreate
```

===== ? VALIDATION & HEALTH CHECKS =====

Quick health check

```
.\validate-platform.ps1
```

Detailed health check with connectivity tests

```
.\validate-platform.ps1 -Detailed
```

Continuous monitoring (30 second intervals)

```
.\validate-platform.ps1 -Continuous -Interval 30
```

Check all container status

```
docker ps -a
```

Check specific service

```
docker ps --filter "name=ollama"
```

? MONITORING COMMANDS

View container logs (live)

```
docker logs -f <container-name>
```

View last 100 lines

```
docker logs --tail 100 <container-name>
```

View logs with timestamps

```
docker logs -t <container-name>
```

Resource usage (real-time)

```
docker stats
```

Resource usage (snapshot)

```
docker stats --no-stream
```

Container inspection

```
docker inspect <container-name>
```

=====

? OLLAMA AI MODEL COMMANDS

=====

List installed models

```
docker exec ollama ollama list
```

Pull new model

```
docker exec ollama ollama pull <model-name>
```

Examples:

```
docker exec ollama ollama pull deepseek-coder:6.7b-instruct-q4_K_M docker exec ollama ollama  
pull qwen2.5-coder:7b-instruct-q4_K_M docker exec ollama ollama pull  
codellama:13b-instruct-q4_K_M
```

Run model interactively

```
docker exec -it ollama ollama run deepseek-coder:33b-instruct-q4_K_M
```

Test model with prompt

```
docker exec ollama ollama run deepseek-coder:33b-instruct-q4_K_M "Write a Python function to  
calculate factorial"
```

Remove model

```
docker exec ollama ollama rm <model-name>
```

Show model info

```
docker exec ollama ollama show deepseek-coder:33b-instruct-q4_K_M
```

? DATABASE COMMANDS

PostgreSQL

Connect to PostgreSQL

```
docker exec -it postgres psql -U modernization -d legacy_modernization
```

Run SQL query

```
docker exec postgres psql -U modernization -d legacy_modernization -c "SELECT version();"
```

Backup database

```
docker exec postgres pg_dump -U modernization legacy_modernization > backup.sql
```

Restore database

```
cat backup.sql | docker exec -i postgres psql -U modernization -d legacy_modernization
```

Redis

Connect to Redis

```
docker exec -it redis redis-cli
```

Get all keys

```
docker exec redis redis-cli KEYS ''
```

Get specific key

```
docker exec redis redis-cli GET <key>
```

Flush all data (CAUTION!)

```
docker exec redis redis-cli FLUSHALL
```

?? STORAGE COMMANDS

MinIO

Access MinIO console: <http://localhost:9001>
Credentials: minioadmin / minioadmin123

ChromaDB

Access ChromaDB API: <http://localhost:8000>
Check health

```
curl http://localhost:8000/api/v1/heartbeat
```

? CONTAINER MANAGEMENT

Start all containers

```
docker start $(docker ps -aq)
```

Stop all containers

```
docker stop $(docker ps -aq)
```

Restart specific container

```
docker restart <container-name>
```

Remove container

```
docker rm <container-name>
```

Remove container (force)

```
docker rm -f <container-name>
```

Recreate container

```
docker rm -f <container-name> .\rebuild-platform.ps1
```

? VOLUME MANAGEMENT

List volumes

```
docker volume ls
```

Inspect volume

```
docker volume inspect <volume-name>
```

Remove unused volumes

```
docker volume prune -f
```

Remove specific volume

```
docker volume rm <volume-name>
```

Backup volume

```
docker run --rm -v <volume-name>:/data -v ${PWD}:/backup alpine tar czf /backup/volume-backup.tar.gz -C /data .
```

Restore volume

```
docker run --rm -v <volume-name>:/data -v ${PWD}:/backup alpine tar xzf /backup/volume-backup.tar.gz -C /data
```

? NETWORK COMMANDS

List networks

```
docker network ls
```

Inspect network

```
docker network inspect modernization-network
```

Connect container to network

```
docker network connect <network-name> <container-name>
```

Disconnect container from network

```
docker network disconnect <network-name> <container-name>
```

Remove network

```
docker network rm <network-name>
```

? CLEANUP COMMANDS

Stop all containers

```
docker stop $(docker ps -aq)
```

Remove all containers

```
docker rm $(docker ps -aq)
```

Remove unused images

```
docker image prune -a -f
```

Remove unused volumes

```
docker volume prune -f
```

Remove unused networks

```
docker network prune -f
```

Complete cleanup (CAUTION: removes everything)

```
docker system prune -a -f --volumes
```

? TROUBLESHOOTING COMMANDS

Check Docker Desktop status

docker info

Check WSL2 integration

wsl -l -v

Check GPU availability

docker run --rm --gpus all nvidia/cuda:12.2.0-base-ubuntu22.04 nvidia-smi

Check port usage

netstat -ano | findstr :<port>

Test connectivity to service

Test-NetConnection -ComputerName localhost -Port <port>

Check container IP address

docker inspect -f '{{range .NetworkSettings.Networks}}{{.IPAddress}}{{end}}' <container-name>

Execute command in container

docker exec <container-name> <command>

Get shell in container

docker exec -it <container-name> /bin/sh docker exec -it <container-name> /bin/bash

? PROMETHEUS QUERIES

Access Prometheus: <http://localhost:9090>

Example queries:

Ollama request rate

```
rate(ollama_request_total[5m])
```

Container CPU usage

```
rate(container_cpu_usage_seconds_total[5m]) * 100
```

Container memory usage (GB)

```
container_memory_usage_bytes / 1024 / 1024 / 1024
```

GPU utilization

```
nvidia_gpu_utilization
```

Disk usage

```
(node_filesystem_size_bytes - node_filesystem_free_bytes) / node_filesystem_size_bytes * 100
```

? GRAFANA COMMANDS

Access Grafana: <http://localhost:3000>

Default credentials: admin / admin123

Add Prometheus data source:

Configuration > Data Sources > Add data source >

URL: `http://prometheus:9090`

Add Loki data source:

Configuration > Data Sources > Add data source >

URL: `http://loki:3100`

? SECURITY COMMANDS

View container environment variables

```
docker exec <container-name> env
```

Check file permissions in container

```
docker exec <container-name> ls -la /path
```

Update password (example: Grafana)

```
docker exec grafana grafana-cli admin reset-admin-password <new-password>
```

? LOGGING COMMANDS

View logs from all containers

```
docker-compose logs -f
```

View logs from specific service

```
docker-compose logs -f <service-name>
```

Export logs to file

```
docker logs <container-name> > container-logs.txt 2>&1
```

Follow logs with grep filter

```
docker logs -f <container-name> | grep ERROR
```

? EMERGENCY COMMANDS

Complete platform restart

```
docker restart $(docker ps -aq)
```

Force stop all containers

```
docker kill $(docker ps -aq)
```

Rebuild everything from scratch

```
docker stop $(docker ps -aq) docker rm $(docker ps -aq) docker volume rm $(docker volume ls -q)  
docker network rm modernization-network monitoring-network .\rebuild-platform.ps1
```

Restart Docker Desktop

Stop-Service docker

Start-Service docker

Or restart via Docker Desktop UI

? SUPPORT COMMANDS

Generate system report

```
docker info > system-report.txt docker ps -a >> system-report.txt docker stats --no-stream >> system-report.txt docker volume ls >> system-report.txt docker network ls >> system-report.txt
```

Export configuration

```
docker-compose config > current-config.yml
```

Check version information

```
docker version docker-compose version
```

END OF QUICK REFERENCE

SOURCE: TROUBLESHOOTING.md

? Troubleshooting Guide

Common Issues and Solutions

? --

? Docker Desktop Issues

Issue: Docker Desktop Won't Start

? *Symptoms:**

- ? Docker Desktop stuck at "Starting..."
- ? Error: "Docker Desktop is not running"

? *Solutions:**

1. **Restart Docker Desktop**: “powershell # Via PowerShell (as Administrator) Stop-Service docker Start-Service docker

Or restart via Docker Desktop UI “

2. **Check WSL2 Status**: “powershell wsl -l -v # Ensure WSL2 is running wsl --update “

3. **Reset Docker Desktop**: - Open Docker Desktop - Settings > Trouble shoot > Reset to factory defaults - **WARNING**: This will delete all containers, images, and volumes

4. **Check Windows Services**: - Open Services (services.msc) - Ensure "Docker Desktop Service" is running - Start type should be "Automatic"

? --

? Port Already in Use

Issue: Port Conflict When Starting Containers

? *Symptoms:**

- ? Error: "Bind for 0.0.0.0:5432 failed: port is already allocated"
- ? Container fails to start

? *Solutions:**

1. **Identify Process Using Port**: “powershell # Find process ID netstat -ano | findstr :<port>

Example for PostgreSQL (port 5432) netstat -ano | findstr :5432 “

2. **Kill Process**: “powershell # Kill process by PID Stop-Process -Id <PID> -Force

Example Stop-Process -Id 12345 -Force “

3. **Change Port in Configuration**: - Edit docker-compose.yml - Change port mapping (e.g., "5433:5432" instead of "5432:5432")

? --

? Container Issues

Issue: Container Keeps Restarting

? *Symptoms:**

- ? Container status shows "Restarting"
- ? Container exits immediately after start

? *Solutions:**

1. **Check Container Logs**: “powershell docker logs <container-name>”
2. **Inspect Container**: “powershell docker inspect <container-name>”
3. **Common Causes**: - **Configuration Error**: Check environment variables - **Port Conflict**: Ensure port is available - **Volume Mount Issue**: Verify volume paths exist - **Insufficient Resources**: Check RAM/CPU allocation
4. **Fix and Restart**: “powershell # Remove problematic container docker rm -f <container-name>

Rebuild .\rebuild-platform.ps1 “

Issue: Container is Unhealthy

? *Symptoms:**

- ? Docker ps shows "unhealthy" status
- ? Health check failing

? *Solutions:**

1. **Check Health Check Logs**: “powershell docker inspect --format='{{json .State.Health}}' <container-name>”
2. **Manual Health Check**: “powershell # Example for Ollama curl http://localhost:11434/

```
# Example for ChromaDB curl http://localhost:8000/api/v1/heartbeat ""
```

3. **Restart Container**: “powershell docker restart <container-name> ”

? --

? Ollama Issues

Issue: GPU Not Detected by Ollama

? *Symptoms:**

? Ollama running on CPU only

? Slow model inference

? nvidia-smi not working in container

? *Solutions:**

1. **Verify Docker GPU Support**: “powershell docker run --rm --gpus all nvidia/cuda:12.2.0-base-ubuntu22.04 nvidia-smi ”

2. **Check Ollama Container GPU Access**: “powershell docker exec ollama nvidia-smi ”

3. **Verify Docker Desktop GPU Settings**: - Docker Desktop > Settings > Resources - Ensure "Use GPU" is enabled - Restart Docker Desktop

4. **Recreate Ollama Container**: “powershell docker rm -f ollama .\rebuild-platform.ps1 ”

Issue: Model Download Fails

? *Symptoms:**

? Error: "pull model manifest"

? Download interrupted

? Insufficient disk space

? *Solutions:**

1. **Check Disk Space**: “powershell # Check available space Get-PSDrive C ”

2. **Retry Download**: “powershell docker exec ollama ollama pull <model-name> ”

3. **Download Smaller Model**: “powershell # Use 7B instead of 33B docker exec ollama ollama pull deepseek-coder:6.7b-instruct-q4_K_M “

4. **Clear Failed Downloads**: “powershell docker exec ollama ollama rm <model-name> “

Issue: Model Inference is Slow

? *Symptoms:**

? Long response times

? High CPU usage

? Out of memory errors

? *Solutions:**

1. **Check GPU Utilization**: “powershell nvidia-smi “

2. **Use Smaller Model or Lower Quantization**: - Q4_K_M (medium quality, faster) - Q2_K (lower quality, fastest)

3. **Reduce Concurrent Requests**: - Limit parallel inference calls - Implement request queuing

4. **Increase GPU Memory**: - Close other GPU-intensive applications - Use model with smaller parameters

? --

? Database Issues

Issue: PostgreSQL Connection Refused

? *Symptoms:**

? Error: "could not connect to server"

? Connection timeout

? *Solutions:**

1. **Verify Container is Running**: “powershell docker ps --filter "name=postgres" “

2. **Check Logs**: “powershell docker logs postgres “

3. **Test Connection**: “powershell docker exec postgres psql -U modernization -d

```
legacy_modernization -c "SELECT version();"
```

4. **Verify Port is Open**: “powershell Test-NetConnection -ComputerName localhost -Port 5432”

Issue: Redis Connection Issues

? *Symptoms:**

? Error: "Could not connect to Redis"

? Commands timing out

? *Solutions:**

1. **Test Redis**: “powershell docker exec redis redis-cli PING # Should return PONG”

2. **Check Redis Logs**: “powershell docker logs redis”

3. **Restart Redis**: “powershell docker restart redis”

? --

? Monitoring Issues

Issue: Prometheus Targets Down

? *Symptoms:**

? Targets showing as "DOWN" in Prometheus UI

? No metrics being collected

? *Solutions:**

1. **Check Target Accessibility**: “powershell # Test Ollama metrics curl http://localhost:11434/metrics”

2. **Verify Network Connectivity**: “powershell docker exec prometheus wget -O http://ollama:11434/metrics”

3. **Check Prometheus Configuration**: “powershell docker exec prometheus cat /etc/prometheus/prometheus.yml”

4. **Reload Prometheus Configuration**: “powershell # Via API Invoke-WebRequest -Method POST

<http://localhost:9090/-/reload>

Or restart docker restart prometheus ““

Issue: Grafana Dashboard Shows No Data

? *Symptoms:**

? Empty graphs

? "No data" message

? *Solutions:**

1. **Verify Data Source Connection**: - Grafana > Configuration > Data Sources - Test connection to Prometheus/Loki
2. **Check Query Syntax**: - Verify PromQL/LogQL queries - Test queries in Prometheus UI first
3. **Verify Time Range**: - Check dashboard time picker - Ensure data exists for selected time range
4. **Check Prometheus Scrape Status**: - Visit <http://localhost:9090/targets> - Verify all targets are UP

? --

?? Storage Issues

Issue: Volume Mount Fails

? *Symptoms:**

? Container can't access volume

? Permission denied errors

? Data not persisting

? *Solutions:**

1. **Check Volume Exists**: “powershell docker volume ls | Select-String <volume-name> ““
2. **Recreate Volume**: “powershell docker volume rm <volume-name> docker volume create <volume-name> ““

3. **Check Volume Permissions**: “powershell docker exec <container-name> ls -la /path/to/mount”

Issue: Disk Space Full

? *Symptoms:**

? "No space left on device"

? Containers failing to start

? Unable to download models

? *Solutions:**

1. **Check Disk Usage**: “powershell docker system df”

2. **Clean Up Docker**: “powershell # Remove unused images docker image prune -a -f

Remove unused volumes (CAUTION) docker volume prune -f

Complete cleanup docker system prune -a -f --volumes”

3. **Move Docker Data Location**: - Docker Desktop > Settings > Resources > Advanced - Change "Disk image location"

? --

? Network Issues

Issue: Containers Can't Communicate

? *Symptoms:**

? Services can't reach each other

? Network timeout errors

? *Solutions:**

1. **Verify Network Exists**: “powershell docker network ls | Select-String modernization”

2. **Check Container Network Connections**: “powershell docker inspect <container-name> --format='{{json .NetworkSettings.Networks}}'”

3. **Reconnect Container to Network**: “powershell docker network connect modernization-

network <container-name> “

4. **Recreate Network**: “powershell docker network rm modernization-network docker network create modernization-network “

? --

? Authentication Issues

Issue: Can't Connect to MinIO Console

? *Symptoms:**

? Login fails

? "Access Denied" error

? *Solutions:**

1. **Verify Credentials**: - Username: ‘minioadmin’ - Password: ‘minioadmin123’

2. **Check Environment Variables**: “powershell docker exec minio env | Select-String MINIO_ROOT “

3. **Reset MinIO**: “powershell docker rm -f minio docker volume rm minio-data .\rebuild-platform.ps1 “

Issue: Grafana Login Issues

? *Symptoms:**

? Default credentials not working

? Locked out of Grafana

? *Solutions:**

1. **Reset Admin Password**: “powershell docker exec grafana grafana-cli admin reset-admin-password newpassword123 “

2. **Check Environment Variables**: “powershell docker exec grafana env | Select-String GF_SECURITY “

? --

? Performance Issues

Issue: High CPU Usage

? *Symptoms:**

? System running slow

? Containers using 100% CPU

? *Solutions:**

1. **Identify Resource Hog**: “powershell docker stats --no-stream “

2. **Limit Container Resources**: “yaml # In docker-compose.yml services: ollama: deploy: resources: limits: cpus: '8' memory: 32G “

3. **Reduce Concurrent Operations**: - Limit parallel model inference - Reduce scrape frequency in Prometheus

Issue: High Memory Usage

? *Symptoms:**

? System running out of RAM

? OOM (Out of Memory) errors

? *Solutions:**

1. **Check Memory Usage**: “powershell docker stats --no-stream --format "table {{.Container}}\t{{.MemUsage}}" “

2. **Increase Docker Memory Allocation**: - Docker Desktop > Settings > Resources - Increase Memory slider

3. **Use Smaller Models**: - Switch to 7B models instead of 33B - Use higher quantization (Q2_K instead of Q4_K_M)

? --

? Configuration Issues

Issue: Changes Not Taking Effect

? *Symptoms:**

? Configuration updates ignored

? Old settings persist

? *Solutions:**

1. **Reload Configuration**: “powershell # For Prometheus Invoke-WebRequest -Method POST http://localhost:9090/-/reload ““

2. **Recreate Containers**: “powershell docker-compose down docker-compose up -d --force-recreate ““

3. **Clear Docker Cache**: “powershell docker-compose build --no-cache docker-compose up -d ““

? --

? Emergency Recovery

Complete Platform Reset

If nothing else works, perform a complete reset:

```
# 1. Stop everything
docker stop $(docker ps -aq)

# 2. Remove all containers
docker rm $(docker ps -aq)

# 3. Remove all volumes (DATA LOSS!)
docker volume prune -f

# 4. Remove all networks
docker network prune -f

# 5. Remove all images (optional)
docker image prune -a -f

# 6. Restart Docker Desktop
# - Docker Desktop > Restart

# 7. Rebuild platform
.\rebuild-platform.ps1
```

? --

? Getting Help

Collect Diagnostic Information

```
# Generate system report
docker info > system-report.txt
docker ps -a >> system-report.txt
docker stats --no-stream >> system-report.txt
docker volume ls >> system-report.txt
docker network ls >> system-report.txt
.\validate-platform.ps1 -Detailed >> system-report.txt
```

Check Service Logs

```
# Export logs from all containers
docker logs postgres > logs/postgres.log 2>&1
docker logs redis > logs/redis.log 2>&1
docker logs ollama > logs/ollama.log 2>&1
docker logs prometheus > logs/prometheus.log 2>&1
```

? --

? Prevention Tips

1. ****Regular Health Checks**:** “powershell .\validate-platform.ps1 -Continuous -Interval 300”
2. ****Monitor Resource Usage**:** - Keep an eye on CPU/RAM/Disk usage - Set up alerts in Prometheus
3. ****Regular Backups**:** - Backup volumes regularly - Export important configurations
4. ****Keep Docker Updated**:** - Update Docker Desktop regularly - Update container images periodically
5. ****Documentation**:** - Document custom configurations - Keep track of changes

? --

? *Last Updated*: 2026-01-16

SOURCE: RUNNER-SETUP.md

GitLab Runner with Docker-in-Docker (DinD) Setup

Complete guide for configuring a GitLab runner capable of building and pushing Docker images to Nexus registry.

Quick Start

```
cd files  
.\\setup-gitlab-runner-dind.ps1 -GitLabUrl "http://gitlab-server" -RunnerName "docke
```

The script will: 1. Verify Docker is running 2. Create the ‘modernization-network’ bridge (if needed) 3. Generate runner configuration with DinD support 4. Optionally start the runner container

Architecture

```
GitLab CI/CD Pipeline (.gitlab-ci.yml)  
?  
  docker_build job  
    ?  
  Runner (docker executor)  
    ?  
  DinD Service (docker:24-dind)  
    ?  
  Docker Socket Sharing (/var/run/docker.sock)  
    ?  
Build & Push to Nexus Registry
```

Configuration Details

Runner Configuration (‘runner-config.toml’)

Key settings for Docker-in-Docker:

```
[ [runners]]  
  name = "docker-dind-runner"  
  executor = "docker"
```

```

[runners.docker]
  image = "docker:24-cli"                      # Lightweight Docker client
  privileged = true                            # Required for Docker operations
  volumes = [
    "/cache",
    "/var/run/docker.sock:/var/run/docker.sock"  # Access host Docker daemon
  ]
  shm_size = 2147483648                         # 2GB for builds
  network_mode = "modernization-network"

[[runners.docker.services]]
  name = "docker:24-dind"                      # Docker-in-Docker service
  alias = "docker"
  command = [ "--tls=false" ]                   # Disable TLS for reliability

```

CI/CD Pipeline Configuration

In your '.gitlab-ci.yml':

```

docker_build:
  stage: docker
  image: docker:24-cli
  services:
    - name: docker:24-dind
      alias: docker
      command: [ "--tls=false" ]

variables:
  DOCKER_TLS_CERTDIR: ""                      # Required for DinD
  DOCKER_HOST: tcp://docker:2375              # DinD endpoint
  DOCKER_DRIVER: overlay2
  FF_NETWORK_PER_BUILD: "true"                 # Per-build network isolation
  FF_DOCKER_LAYER_CACHING: "true"              # Cache Docker layers
  DOCKER_REGISTRY: nexus-docker:5001
  DOCKER_IMAGE_NAME: java-app

before_script:
  - sleep 5                                     # Wait for DinD service
  - docker info                                  # Verify connectivity
  - docker login -u $DOCKER_REGISTRY_USER -p $DOCKER_REGISTRY_PASSWORD $DOCKER_REGISTRY

script:
  - docker build -t $DOCKER_REGISTRY/$DOCKER_IMAGE_NAME:latest .
  - docker push $DOCKER_REGISTRY/$DOCKER_IMAGE_NAME:latest

```

Manual Runner Registration

If you prefer manual setup instead of the PowerShell script:

1. Create Configuration File

Create ‘`~/.gitlab-runner/config.toml`’:

```
concurrent = 3
check_interval = 0

[session_server]
  session_timeout = 1800

[[runners]]
  name = "docker-dind-runner"
  url = "http://gitlab-server"
  token = "YOUR_REGISTRATION_TOKEN"
  executor = "docker"

[runners.docker]
  image = "docker:24-cli"
  privileged = true
  volumes = [ "/cache", "/var/run/docker.sock:/var/run/docker.sock" ]
  shm_size = 2147483648
  network_mode = "modernization-network"

[[runners.docker.services]]
  name = "docker:24-dind"
  alias = "docker"
  command = [ "--tls=false" ]
```

2. Run the Runner Container

```
docker run -d \
  --name gitlab-runner \
  --restart unless-stopped \
  --network modernization-network \
  -v ~/.gitlab-runner/config.toml:/etc/gitlab-runner/config.toml:ro \
  -v /var/run/docker.sock:/var/run/docker.sock \
  gitlab/gitlab-runner:latest
```

3. Verify the Runner is Running

```
docker logs -f gitlab-runner
```

GitLab CI/CD Variables Setup

Set these in GitLab (Settings > CI/CD > Variables):

Variable	Value	Masked
'DOCKER_REGISTRY_USER'	Nexus username	No
'DOCKER_REGISTRY_PASSWORD'	Nexus password	**Yes**
'DOCKER_REGISTRY'	'nexus-docker:5001'	No
'DOCKER_IMAGE_NAME'	'java-app'	No

Network Configuration

The runner must be on the same Docker network as Nexus:

```
# Verify network exists
docker network ls | grep modernization-network

# Create if missing
docker network create modernization-network
```

All services communicate by container name: - 'nexus-docker:5001' ? Nexus Docker registry - 'docker:2375' ? DinD service endpoint

Troubleshooting

Runner Won't Connect to GitLab

```
# Check runner logs
docker logs gitlab-runner

# Verify GitLab is reachable
docker exec gitlab-runner ping gitlab-server

# Test from runner container
docker run --rm --network modernization-network curlimages/curl curl http://gitlab-
```

Docker Commands Fail in Pipeline

```
# Check DinD service connectivity
docker exec gitlab-runner docker info

# Verify socket binding
docker inspect gitlab-runner | grep -A 10 Mounts
```

Image Push Fails to Nexus

```
# Verify Nexus is accessible
docker run --rm --network modernization-network curlimages/curl curl http://nexus-ci:8081/api/v3/projects

# Check registry credentials in pipeline logs
# (ensure DOCKER_REGISTRY_PASSWORD is masked in CI/CD variables)

# Test login locally
docker login -u admin -p admin123 nexus-docker:5001
```

Out of Disk Space During Builds

The ‘shm_size = 2147483648’ (2GB) might need adjustment:
- Increase for large Docker image builds
- Decrease if system RAM is limited

Update in ‘runner-config.toml’ and restart runner:

```
docker restart gitlab-runner
```

Maintenance

View Active Jobs

```
docker exec gitlab-runner gitlab-runner --debug list
```

Update Runner Version

```
docker pull gitlab/gitlab-runner:latest
docker restart gitlab-runner
```

Rotate Runner Token

1. In GitLab: Settings > CI/CD > Runners > Click runner > Rotate token
2. Update ‘runner-config.toml’ with new token
3. Restart runner: ‘docker restart gitlab-runner’

Performance Tips

1. **Cache Docker Layers**: Enable ‘FF_DOCKER_LAYER_CACHING’ to reuse build layers
2. **Parallel Builds**: Set ‘concurrent = 3’ (or higher) for multiple simultaneous jobs
3. **Shared Memory**: Increase ‘shm_size’ if experiencing OOM errors
4. **Network**: Use

```
'network_mode = "modernization-network" for service communication
```

Security Considerations

?? **Production Deployment**: - Use TLS with proper certificates (set 'tls_verify = true') - Don't share '/var/run/docker.sock' in production (security risk) - Use dedicated runner host with proper isolation - Rotate credentials regularly - Use masked CI/CD variables for secrets - Restrict runner to specific projects/tags

For production, consider using Kubernetes executor with isolated pod security policies instead of Docker executor.

References

- ? [GitLab Runner Documentation](<https://docs.gitlab.com/runner/>)
- ? [Docker-in-Docker Executor](<https://docs.gitlab.com/runner/executors/docker.html#docker-in-docker>)
- ? [Docker Service Variables](<https://docs.gitlab.com/ee/ci/services/>)
- ? [Nexus Docker Repository](<https://help.sonatype.com/repomanager3/formats/docker-registry>)

SOURCE: RUNNER-SETUP-SUMMARY.md

GitLab Runner DinD Configuration - Complete Setup

What Was Configured

Your GitLab runner is now configured with **Docker-in-Docker (DinD)** capability for building and pushing Docker images to Nexus registry.

Files Modified/Created

1. Runner Configuration

- ? *java-pipeline/runner-config.toml*
- ? Updated concurrent jobs from 1 to 3
- ? Added DinD service configuration with 'docker:24-dind'

- ? Enabled Docker socket sharing ('/var/run/docker.sock')
- ? Configured 'modernization-network' for service communication
- ? Set 2GB shared memory for Docker operations

2. CI/CD Pipeline

- ? *java-pipeline/.gitlab-ci.yml**
- ? Enhanced Docker build job with DinD variables
- ? Added environment variables for Docker registry
- ? Configured DinD service with TLS disabled (for reliability)
- ? Added layer caching flags for faster builds

3. Setup Script (NEW)

- ? *files/setup-gitlab-runner-dind.ps1**
- ? PowerShell script for automatic runner setup
- ? Creates runner configuration with DinD enabled
- ? Starts GitLab runner container
- ? Provides comprehensive setup guidance

4. Documentation (NEW)

- ? *files/RUNNER-SETUP.md**
- ? Complete runner setup guide
- ? Configuration explanations
- ? Troubleshooting section
- ? Security considerations
- ? Performance optimization tips

Quick Setup Steps

Step 1: Run the Setup Script

```
cd files
.\setup-gitlab-runner-dind.ps1 -GitLabUrl "http://gitlab-server" -RunnerName "docke
```

Step 2: Configure GitLab CI/CD Variables

In GitLab UI (Settings > CI/CD > Variables): - 'DOCKER_REGISTRY_USER': Your Nexus username - 'DOCKER_REGISTRY_PASSWORD': Your Nexus password (mark as **Masked**) - 'DOCKER_REGISTRY': 'nexus-docker:5001' - 'DOCKER_IMAGE_NAME': 'java-app'

Step 3: Verify Runner is Running

```
docker logs -f gitlab-runner
```

Key Configuration Points

DinD Service Setup

```
services:
  - name: docker:24-dind
    alias: docker
    command: ["--tls=false"] # Reliability without TLS in dev
```

Docker Socket Sharing

```
volumes = [
  "/cache",
  "/var/run/docker.sock:/var/run/docker.sock" # Access host Docker
]
```

Environment Variables

```
DOCKER_TLS_CERTDIR: ""                      # Required for DinD
DOCKER_HOST: tcp://docker:2375                # DinD endpoint
DOCKER_DRIVER: overlay2                       # Efficient storage
FF_DOCKER_LAYER_CACHING: "true"               # Reuse build layers
```

How It Works

1. GitLab CI/CD triggers `docker_build` job
?
2. Job runs in `docker:24-cli` container
?
3. DinD service (`docker:24-dind`) provides Docker daemon
?
4. Job executes `docker build/push` commands
?
5. Image gets pushed to `nexus-docker:5001` registry

Docker Build Commands in Your Pipeline

Example in `.gitlab-ci.yml`:

```
yaml script: - echo "$DOCKER_REGISTRY_PASSWORD" | docker login -u "$DOCKER_REGISTRY_USER" --password-stdin $DOCKER_REGISTRY - docker build -t $DOCKER_REGISTRY/$DOCKER_IMAGE_NAME:${CI_COMMIT_SHA:0:8} . - docker push $DOCKER_REGISTRY/$DOCKER_IMAGE_NAME:${CI_COMMIT_SHA:0:8}
```

Networking

- ? All components use ‘modernization-network’ bridge
- ? Service discovery by container name:
- ? ‘docker:2375’ ? DinD daemon
- ? ‘nexus-docker:5001’ ? Nexus registry
- ? ‘gitlab-server’ ? GitLab instance

Validation

Check runner status: ““bash # View runner logs docker logs gitlab-runner

Test Docker connectivity

docker exec gitlab-runner docker ps

List active jobs

docker exec gitlab-runner gitlab-runner --debug list ““

Next Steps

1. Register runner with GitLab if not done automatically
2. Set CI/CD variables for Nexus authentication
3. Trigger a test pipeline to verify Docker build works
4. Monitor logs: ‘docker logs -f gitlab-runner’

Troubleshooting Reference

See files/RUNNER-SETUP.md for detailed troubleshooting:
- Runner connection issues
- Docker daemon connectivity
- Nexus registry authentication
- Disk space and memory issues

Security Notes

- ? **Development**: Current setup is fine for local development ?? **Production**: - Enable TLS for DinD - Don’t share docker.sock (use Kubernetes executor instead) - Use proper registry authentication - Rotate credentials regularly - Add network policies

For production, consider Kubernetes executor with pod security policies.

SOURCE: RUNNER-QUICK-REFERENCE.md

GitLab Runner DinD - Quick Reference Card

? Files Modified/Created

File	Change	Purpose
'java-pipeline/runner-config.toml'	?? Modified	DinD runner config with socket
'java-pipeline/.gitlab-ci.yml'	?? Modified	Enhanced Docker build job variables
'files/setup-gitlab-runner-dind.ps1'	? Created	Automated runner setup script
'files/RUNNER-SETUP.md'	? Created	Complete setup guide & troubleshooting
'files/RUNNER-SETUP-SUMMARY.md'	? Created	Configuration overview
'files/gitlab-ci-dind-example.yml'	? Created	Complete pipeline example

? Quick Start

```
# 1. Run setup script
cd files
.\setup-gitlab-runner-dind.ps1

# 2. Enter registration token when prompted (from GitLab)
# 3. Script will start the runner

# 4. Set CI/CD variables in GitLab UI
#     Settings > CI/CD > Variables:
#     - DOCKER_REGISTRY_USER (your Nexus username)
#     - DOCKER_REGISTRY_PASSWORD (masked)
```

? Key Configuration

DinD Service

```
services:
  - name: docker:24-dind
    alias: docker
    command: [ "--tls=false" ]
```

Environment Variables

```
DOCKER_TLS_CERTDIR: ""                      # Required for DinD
DOCKER_HOST: tcp://docker:2375               # DinD endpoint
DOCKER_DRIVER: overlay2                       # Storage driver
```

Runner Settings

```
privileged = true                           # Docker access
volumes = [
    "/var/run/docker.sock:/var/run/docker.sock"  # Socket sharing
]
shm_size = 2147483648                      # 2GB shared memory
```

? Verification Steps

```
# 1. Check runner is running
docker ps | grep gitlab-runner

# 2. View logs
docker logs -f gitlab-runner

# 3. Test Docker in runner
docker exec gitlab-runner docker ps

# 4. Trigger test pipeline in GitLab
# Watch the docker_build job in GitLab UI
```

? GitLab CI/CD Variables Required

```
DOCKER_REGISTRY_USER      = <nexus-username>
DOCKER_REGISTRY_PASSWORD = <nexus-password> [MASKED]
DOCKER_REGISTRY          = nexus-docker:5001
DOCKER_IMAGE_NAME        = java-app
```

? How DinD Works

```
CI Pipeline
?
Docker executor (docker:24-cli)
?
DinD service (docker:24-dind) provides Docker daemon
?
Runner can execute: docker build, docker push, etc.
?
Images pushed to Nexus registry (5001)
```

? Common Commands

```
# View runner logs  
docker logs gitlab-runner  
  
# Restart runner  
docker restart gitlab-runner  
  
# Check Docker daemon in runner  
docker exec gitlab-runner docker info  
  
# List active jobs  
docker exec gitlab-runner gitlab-runner --debug list  
  
# Rotate runner token (after updating in GitLab)  
docker restart gitlab-runner
```

?? Important Notes

1. **TLS Disabled**: Using ‘--tls=false’ for development reliability
2. **Privileged Mode**: Required for Docker operations in CI
3. **Socket Sharing**: Allows accessing host Docker daemon
4. **Network**: Runner must be on ‘modernization-network’
5. **Credentials**: Always mask sensitive variables in GitLab

? Troubleshooting Quick Links

- ? Runner won’t connect: Check GitLab URL and token in config
- ? Docker commands fail: Verify ‘docker info’ works in runner
- ? Image push fails: Check Nexus credentials and registry URL
- ? Out of memory: Increase ‘shm_size’ in config

See ‘files/RUNNER-SETUP.md’ for detailed troubleshooting.

? Full Documentation

- ? **Setup Guide**: ‘files/RUNNER-SETUP.md’
- ? **Configuration Summary**: ‘files/RUNNER-SETUP-SUMMARY.md’
- ? **Pipeline Example**: ‘files/gitlab-ci-dind-example.yml’
- ? **Runner Config**: ‘java-pipeline/runner-config.toml’
- ? **CI/CD Config**: ‘java-pipeline/.gitlab-ci.yml’

? Next Steps

1. ? Run 'setup-gitlab-runner-dind.ps1' to create/start runner
2. ? Get registration token from GitLab (Settings > CI/CD > Runners)
3. ? Set Docker registry credentials in GitLab variables
4. ? Push changes to trigger 'docker_build' job
5. ? Verify images appear in Nexus registry